

Artificial Intelligence

Lecture 8: Bayesian Network II


Xiaojin Gong

2021-04-26

Credits: AI Course in Berkeley

Review

- Bayesian Network
 - Representation
 - Joint Probability
 - Conditional Independence
 - Inference
 - Enumeration
 - Variable Elimination
 - Rejection Sampling
 - Likelihood Weighting
 - Gibbs Sampling




Outline


- Bayesian network
 - Representation
 - Template models
 - Inference
 - Applications

Template Models

- Temporal Modeling
 - State-observation model
 - Hidden Markov model
 - Dynamic Bayesian network
- Object-relational models
 - Plate models




Repetition



Temporal Models

- Distributions over Time
 - Discretize the timeline into a set of time slices
 - X_t : the set of state variables at time $t\Delta$
 - $X_{t:t'} = \{X_t, \dots, X_{t'}\} \quad (t \leq t')$

$$P(X_{0:T}) = P(X_0) \prod_{t=0}^{T-1} P(X_{t+1} | X_{0:t})$$



Markov Chains / Markov Processes

- Markov Assumption:
 - The current state depends on only a finite fixed number of previous states.
 - Such systems are called Markovian.
- First-order Markov chains
 - A dynamic system over the template variable X satisfies the Markov assumption if, for all $t \geq 0$,

$$(X_{t+1} \perp X_{0:t-1} | X_t)$$

- The distribution over time:


$$P(X_{0:T}) = P(X_0) \prod_{t=0}^{T-1} P(X_{t+1} | X_t)$$

- A Markovian dynamic system is stationary (time invariance) if $P(X_{t+1} | X_t)$ is the same for all t .

Markov Chains


- First-order Markov chain

- $(X_{t+1} \perp X_{0:t-1} | X_t)$



- Second-order Markov chain

- $(X_{t+1} \perp X_{0:t-2} | X_t, X_{t-1})$



State-Observation Models


- Two independence assumptions
 - The state variables evolve in a Markovian way

$$(X_{t+1} \perp X_{0:t-1} | X_t)$$

- The observation variables at time t conditionally independent of the entire state sequence given the state variables at time t :

$$(E_t \perp X_{0:t-1}, X_{t+1:\infty} | X_t)$$

$$(E_t \perp E_{1:t-1}, E_{t+1:\infty} | X_t)$$



State-Observation Models

- The transition model:

$$\mathbf{P}(\mathbf{X}_t \mid \mathbf{X}_{0:t-1}) = \mathbf{P}(\mathbf{X}_t \mid \mathbf{X}_{t-1})$$

- The sensor/observation model:


$$\mathbf{P}(\mathbf{E}_t \mid \mathbf{X}_{0:t}, \mathbf{E}_{0:t-1}) = \mathbf{P}(\mathbf{E}_t \mid \mathbf{X}_t)$$

- The initial state model: $P(X_0)$
- Joint distribution:

$$\mathbf{P}(\mathbf{X}_{0:t}, \mathbf{E}_{1:t}) = \mathbf{P}(\mathbf{X}_0) \prod_{i=1}^t \mathbf{P}(\mathbf{X}_i \mid \mathbf{X}_{i-1}) \mathbf{P}(\mathbf{E}_i \mid \mathbf{X}_i)$$

Hidden Markov Models

- HMM is a state-observation model in which the state is described by a single discrete random variable.




- Applications:
 - Speech recognition
 - Part-of-speech tagging

Dynamic Bayesian Networks

- Dynamic Bayesian Network (DBN) represents a temporal probability model in which each slice can have any number of state variables and evidence variables
- A DBN contains:
 - The prior distribution over the state variables $P(X_0)$
 - The transition model $P(X_{t+1}|X_t)$
 - The sensor model $P(E_t|X_t)$

Dynamic Bayesian Networks


- An initial state distribution $P(X_0)$
 - Eg. Monitoring a battery-powered robot moving in the X-Y plane



Dynamic Bayesian Networks

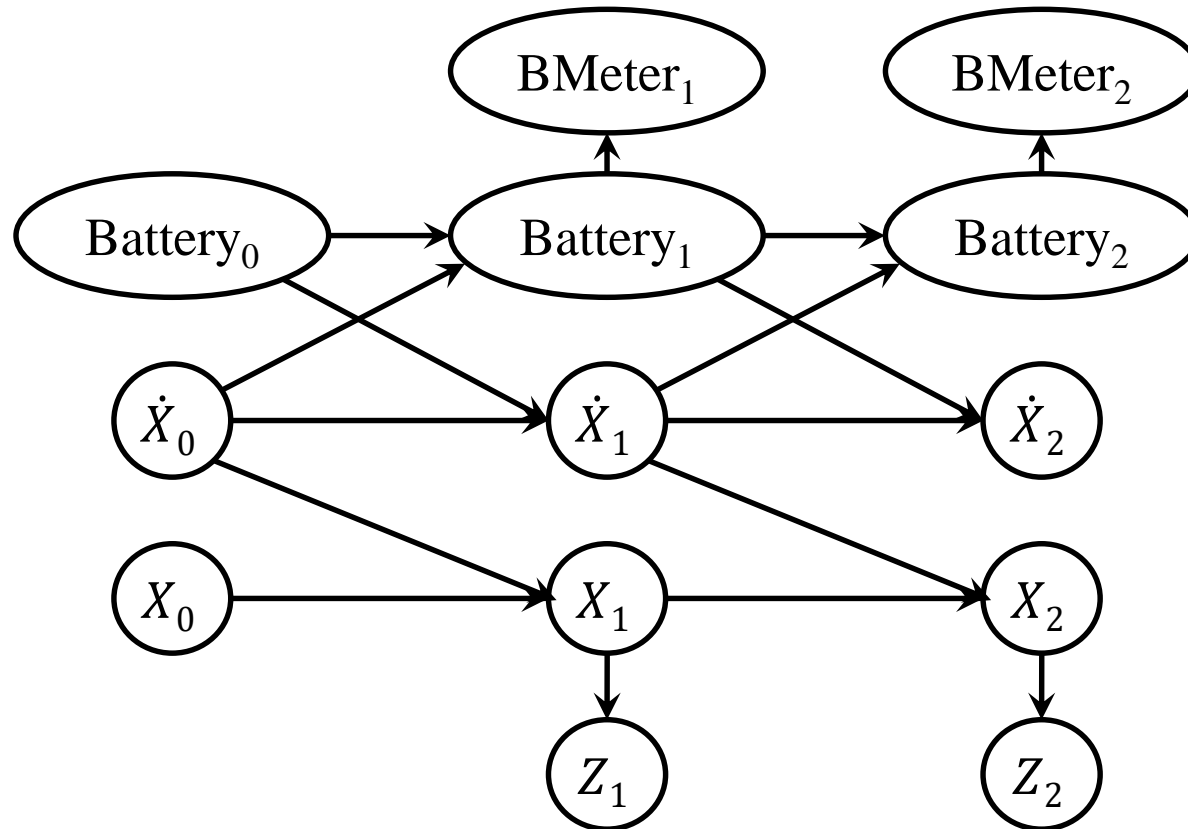
- A transition model $P(X_{t+1}|X_t)$
 - A 2-time-slice Bayesian Network (2TBN) over X^1, \dots, X^n is specified as a BN fragment s.t.
 - The nodes include $X_{t+1}^1, \dots, X_{t+1}^n$ and a subset of X_t^1, \dots, X_t^n .
 - Only the nodes $X_{t+1}^1, \dots, X_{t+1}^n$ have parents and a CPD
 - The 2TBN defines a conditional distribution

$$P(X_{t+1}|X_t) = \prod P(X_{t+1}^i | Pa(X_{t+1}^i))$$



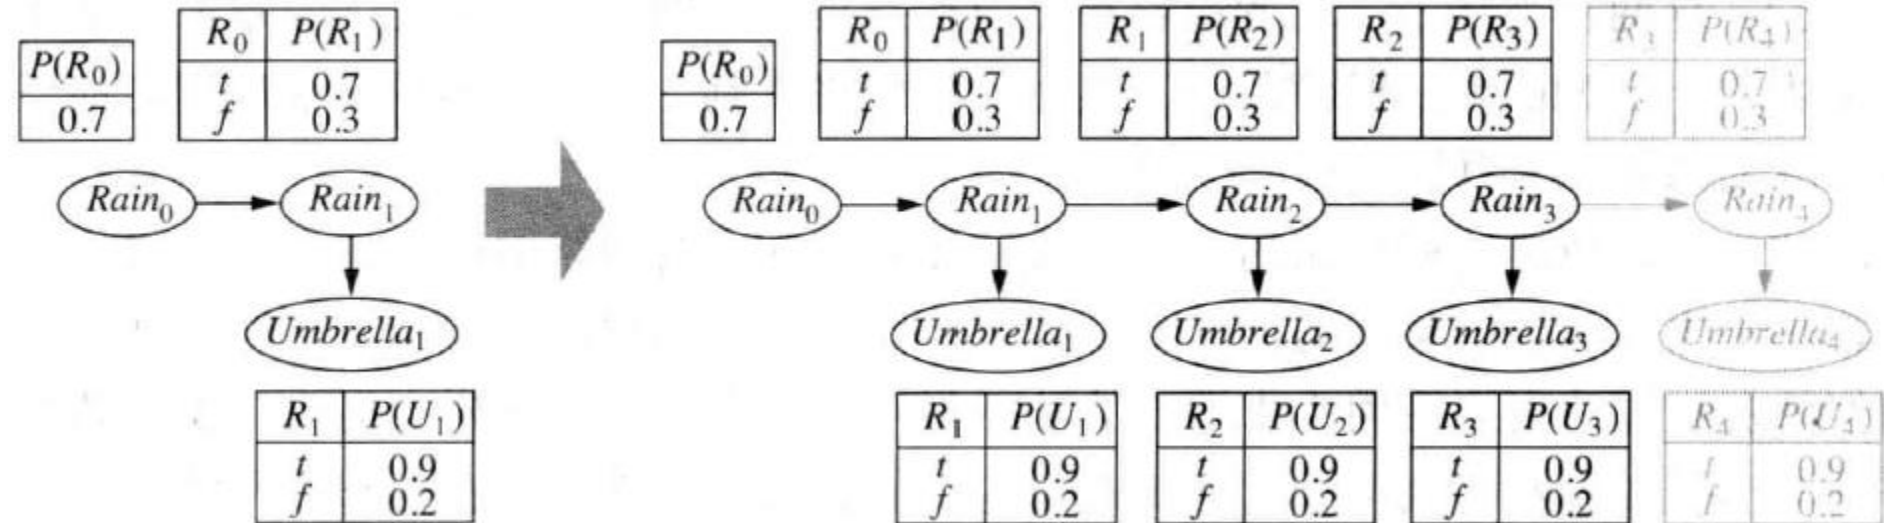
Dynamic Bayesian Networks

- A dynamic Bayesian network (DBN) over X^1, \dots, X^n is defined by
 - A Bayesian network BN_0 over X_0^1, \dots, X_0^n
 - A 2TBN BN_{\rightarrow} over X^1, \dots, X^n




Dynamic Bayesian Networks

- A dynamic Bayesian network (DBN) over X^1, \dots, X^n is defined by
 - A Bayesian network BN_0 over X_0^1, \dots, X_0^n
 - A 2TBN BN_{\rightarrow} over X^1, \dots, X^n



DBN vs. HMM



- Every HMM can be represented as DBN with a single state variable and a single evidence variable.
- Every discrete variable DBN can be represented as an HMM by combining all the state variables into a single state variable



DBN is sparser.

Plate Models

- Plate models are the simplest object-relational models.



Inference in Temporal Models

- Filtering: $P(X_t | e_{1:t})$
 - Computing the **belief state** (the posterior distribution over the most recent state) given all evidence to date.
 - Also called **state estimation**
- Prediction: $P(X_{t+k} | e_{1:t}) \quad k > 0$
 - Computing the posterior distribution over the future state given all evidence to date.
- Smoothing: $P(X_k | e_{1:t}) \quad 0 \leq k < t$
 - Better estimate of past states, essential for learning
- Most likely explanation: $\operatorname{argmax}_{X_{1:t}} P(X_{1:t} | e_{1:t})$
 - Speech recognition
- Learning:
 - Learn the transition and sensor models from observations

Filtering

- Compute $P(X_t | e_{1:t})$
- Recursive estimation
 - Recursively perform prediction and update

$$\begin{aligned} \mathbf{P}(\mathbf{X}_{t+1} | \mathbf{e}_{1:t+1}) &= \mathbf{P}(\mathbf{X}_{t+1} | \mathbf{e}_{1:t}, \mathbf{e}_{t+1}) \\ &= \alpha \mathbf{P}(\mathbf{e}_{t+1} | \mathbf{X}_{t+1}, \mathbf{e}_{1:t}) \mathbf{P}(\mathbf{X}_{t+1} | \mathbf{e}_{1:t}) \\ &= \alpha \underbrace{\mathbf{P}(\mathbf{e}_{t+1} | \mathbf{X}_{t+1})}_{\text{Update}} \underbrace{\mathbf{P}(\mathbf{X}_{t+1} | \mathbf{e}_{1:t})}_{\text{Prediction}} \end{aligned}$$

$$\begin{aligned} \boxed{\mathbf{P}(\mathbf{X}_{t+1} | \mathbf{e}_{1:t+1})} &= \alpha \mathbf{P}(\mathbf{e}_{t+1} | \mathbf{X}_{t+1}) \sum_{\mathbf{x}_t} \mathbf{P}(\mathbf{X}_{t+1} | \mathbf{x}_t, \mathbf{e}_{1:t}) P(\mathbf{x}_t | \mathbf{e}_{1:t}) \\ &= \alpha \underbrace{\mathbf{P}(\mathbf{e}_{t+1} | \mathbf{X}_{t+1})}_{\text{The sensor model}} \sum_{\mathbf{x}_t} \underbrace{\mathbf{P}(\mathbf{X}_{t+1} | \mathbf{x}_t)}_{\text{The transition model}} \boxed{P(\mathbf{x}_t | \mathbf{e}_{1:t})} \end{aligned}$$

Filtering

- View as message passing

$$\mathbf{P}(\mathbf{X}_{t+1} \mid \mathbf{e}_{1:t+1}) = \alpha \mathbf{P}(\mathbf{e}_{t+1} \mid \mathbf{X}_{t+1}) \sum_{\mathbf{x}_t} \mathbf{P}(\mathbf{X}_{t+1} \mid \mathbf{x}_t) P(\mathbf{x}_t \mid \mathbf{e}_{1:t})$$

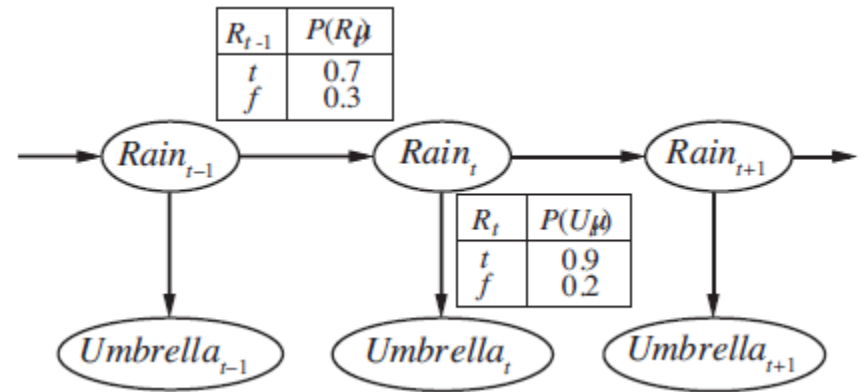
- Consider $\mathbf{P}(\mathbf{X}_t \mid \mathbf{e}_{1:t})$ as a message $\mathbf{f}_{1:t}$
 - Propagated forward along the sequence
 - Modified by each transition
 - Updated by each new observation

$$\mathbf{f}_{1:t+1} = \alpha \text{FORWARD}(\mathbf{f}_{1:t}, \mathbf{e}_{t+1})$$


- Constant time and space for each update

Filtering

- Example:
 - Compute $\mathbf{P}(R_2 \mid u_{1:2})$



$$\begin{aligned} \mathbf{P}(R_2 \mid u_1, u_2) &= \alpha \mathbf{P}(u_2 \mid R_2) \mathbf{P}(R_2 \mid u_1) = \alpha \langle 0.9, 0.2 \rangle \langle 0.627, 0.373 \rangle \\ \mathbf{P}(R_0) &= \langle 0.5, 0.5 \rangle \\ &= \alpha \langle 0.565, 0.075 \rangle \approx \langle 0.883, 0.117 \rangle . \end{aligned}$$



$$\begin{aligned} \mathbf{P}(\mathbf{X}_{t+1} \mid \mathbf{e}_{1:t+1}) &= \alpha \mathbf{P}(\mathbf{e}_{t+1} \mid \mathbf{X}_{t+1}) \mathbf{P}(\mathbf{X}_{t+1} \mid \mathbf{e}_{1:t}) \\ &= \alpha \mathbf{P}(\mathbf{e}_{t+1} \mid \mathbf{X}_{t+1}) \sum_{\mathbf{x}_t} \mathbf{P}(\mathbf{X}_{t+1} \mid \mathbf{x}_t) P(\mathbf{x}_t \mid \mathbf{e}_{1:t}) \end{aligned} \quad 21$$

Prediction

- Compute $P(X_{t+k}|e_{1:t})$
- Viewed as filtering without the addition of new evidence.
- Recursively computed by

$$\boxed{\mathbf{P}(\mathbf{X}_{t+k+1} \mid \mathbf{e}_{1:t})} = \sum_{\mathbf{x}_{t+k}} \underbrace{\mathbf{P}(\mathbf{X}_{t+k+1} \mid \mathbf{x}_{t+k})}_{\text{The transition model}} \boxed{P(\mathbf{x}_{t+k} \mid \mathbf{e}_{1:t})}$$

The transition model

Smoothing

- Compute $P(X_k | e_{1:t})$ $0 \leq k < t$
- Smoothing provides a better estimate of the state than was available at the time
- Learning requires smoothing, rather than filtering
- Forward-backward algorithm

$$\begin{aligned} \mathbf{P}(\mathbf{X}_k | \mathbf{e}_{1:t}) &= \mathbf{P}(\mathbf{X}_k | \mathbf{e}_{1:k}, \mathbf{e}_{k+1:t}) \\ &= \alpha \mathbf{P}(\mathbf{X}_k | \mathbf{e}_{1:k}) \mathbf{P}(\mathbf{e}_{k+1:t} | \mathbf{X}_k, \mathbf{e}_{1:k}) \\ &= \alpha \underbrace{\mathbf{P}(\mathbf{X}_k | \mathbf{e}_{1:k})}_{\substack{\text{Filtering} \\ \text{Forward}}} \underbrace{\mathbf{P}(\mathbf{e}_{k+1:t} | \mathbf{X}_k)}_{\text{Backward}} \\ &= \alpha \mathbf{f}_{1:k} \times \mathbf{b}_{k+1:t} \end{aligned}$$


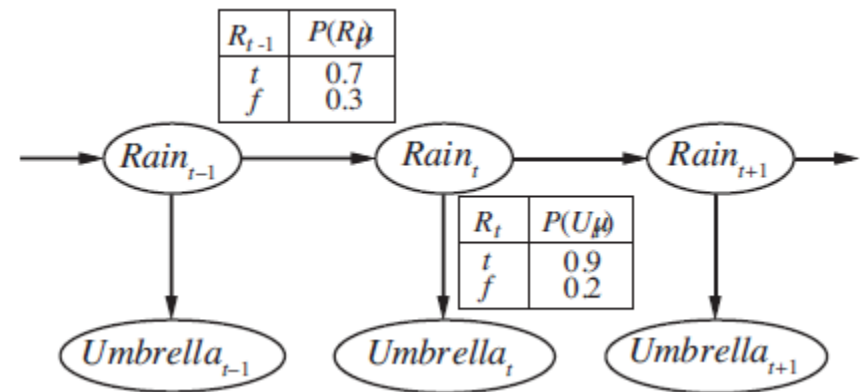
Smoothing

- $\mathbf{P}(\mathbf{X}_k \mid \mathbf{e}_{1:t}) = \alpha \mathbf{P}(\mathbf{X}_k \mid \mathbf{e}_{1:k}) \mathbf{P}(\mathbf{e}_{k+1:t} \mid \mathbf{X}_k)$
 $= \alpha \mathbf{f}_{1:k} \times \mathbf{b}_{k+1:t}$
- Backward message: $\mathbf{b}_{k+1:t} = \mathbf{P}(\mathbf{e}_{k+1:t} \mid \mathbf{X}_k) = \text{BACKWARD}(\mathbf{b}_{k+2:t}, \mathbf{e}_{k+1})$

$$\begin{aligned}
 \boxed{\mathbf{P}(\mathbf{e}_{k+1:t} \mid \mathbf{X}_k)} &= \sum_{\mathbf{x}_{k+1}} \mathbf{P}(\mathbf{e}_{k+1:t} \mid \mathbf{X}_k, \mathbf{x}_{k+1}) \mathbf{P}(\mathbf{x}_{k+1} \mid \mathbf{X}_k) \\
 &= \sum_{\mathbf{x}_{k+1}} P(\mathbf{e}_{k+1:t} \mid \mathbf{x}_{k+1}) \mathbf{P}(\mathbf{x}_{k+1} \mid \mathbf{X}_k) \\
 &= \sum_{\mathbf{x}_{k+1}} P(\mathbf{e}_{k+1}, \mathbf{e}_{k+2:t} \mid \mathbf{x}_{k+1}) \mathbf{P}(\mathbf{x}_{k+1} \mid \mathbf{X}_k) \\
 &= \sum_{\mathbf{x}_{k+1}} \underbrace{P(\mathbf{e}_{k+1} \mid \mathbf{x}_{k+1})}_{\text{Sensor model}} \boxed{P(\mathbf{e}_{k+2:t} \mid \mathbf{x}_{k+1})} \underbrace{\mathbf{P}(\mathbf{x}_{k+1} \mid \mathbf{X}_k)}_{\text{Transition model}}
 \end{aligned}$$

Smoothing

- Example:
 - Compute $\mathbf{P}(R_1 | u_1, u_2)$



$$\mathbf{P}(R_1 | u_1, u_2) = \alpha \langle 0.818, 0.182 \rangle \times \langle 0.69, 0.41 \rangle \approx \langle 0.883, 0.117 \rangle$$

$$\mathbf{P}(\mathbf{X}_k | \mathbf{e}_{1:t}) = \alpha \mathbf{P}(\mathbf{X}_k | \mathbf{e}_{1:k}) \mathbf{P}(\mathbf{e}_{k+1:t} | \mathbf{X}_k) = \alpha \mathbf{f}_{1:k} \times \mathbf{b}_{k+1:t}$$

The Most Likely Sequence

- $\operatorname{argmax}_{X_{1:t}} P(X_{1:t} | e_{1:t})$
- Viterbi algorithm

$$\max_{\mathbf{x}_1 \dots \mathbf{x}_t} \mathbf{P}(\mathbf{x}_1, \dots, \mathbf{x}_t, \mathbf{X}_{t+1} | \mathbf{e}_{1:t+1})$$

$$= \alpha \mathbf{P}(\mathbf{e}_{t+1} | \mathbf{X}_{t+1}) \max_{\mathbf{x}_t} \left(\mathbf{P}(\mathbf{X}_{t+1} | \mathbf{x}_t) \max_{\mathbf{x}_1 \dots \mathbf{x}_{t-1}} P(\mathbf{x}_1, \dots, \mathbf{x}_{t-1}, \mathbf{x}_t | \mathbf{e}_{1:t}) \right)$$

- Identical to filtering if:
 - Replace the forward message $\mathbf{f}_{1:t} = \mathbf{P}(\mathbf{X}_t | \mathbf{e}_{1:t})$ by



$$\mathbf{m}_{1:t} = \max_{\mathbf{x}_1 \dots \mathbf{x}_{t-1}} \mathbf{P}(\mathbf{x}_1, \dots, \mathbf{x}_{t-1}, \mathbf{X}_t | \mathbf{e}_{1:t})$$

- Replace the summation over \mathbf{x}_t by the maximization over \mathbf{x}_t

$$\mathbf{P}(\mathbf{X}_{t+1} | \mathbf{e}_{1:t+1}) = \alpha \mathbf{P}(\mathbf{e}_{t+1} | \mathbf{X}_{t+1}) \sum_{\mathbf{x}_t} \mathbf{P}(\mathbf{X}_{t+1} | \mathbf{x}_t) P(\mathbf{x}_t | \mathbf{e}_{1:t})$$


The Most Likely Sequence

- Example
 - Given the umbrella sequence $[true, true, false, true, true]$
 - Find the weather sequence most likely to explain this




Kalman Filtering

- A dynamic Bayesian network
 - Transition model
 - Sensor model
- } Linear Gaussian Distributions
- The next state is a linear function of the current state, plus some Gaussian noise



Kalman Filtering

- Inference



$$\mathbf{P}(\mathbf{X}_{t+1} | \mathbf{e}_{1:t+1}) = \alpha \mathbf{P}(\mathbf{e}_{t+1} | \mathbf{X}_{t+1}) \sum_{\mathbf{x}_t} \mathbf{P}(\mathbf{X}_{t+1} | \mathbf{x}_t) P(\mathbf{x}_t | \mathbf{e}_{1:t})$$

Kalman Filtering

- Inference

- Predict: (time update $X_t \rightarrow X_{t+1}$)

$$P(X_{t+1}|e_{1:t}) = \int P(X_{t+1}|x_t)P(x_t|e_{1:t})dx_t$$

- Update: (measurement update $e_{t+1} \rightarrow X_{t+1}$)

$$P(X_{t+1}|e_{1:t+1}) = \alpha P(e_{t+1}|X_{t+1})P(X_{t+1}|e_{1:t})$$

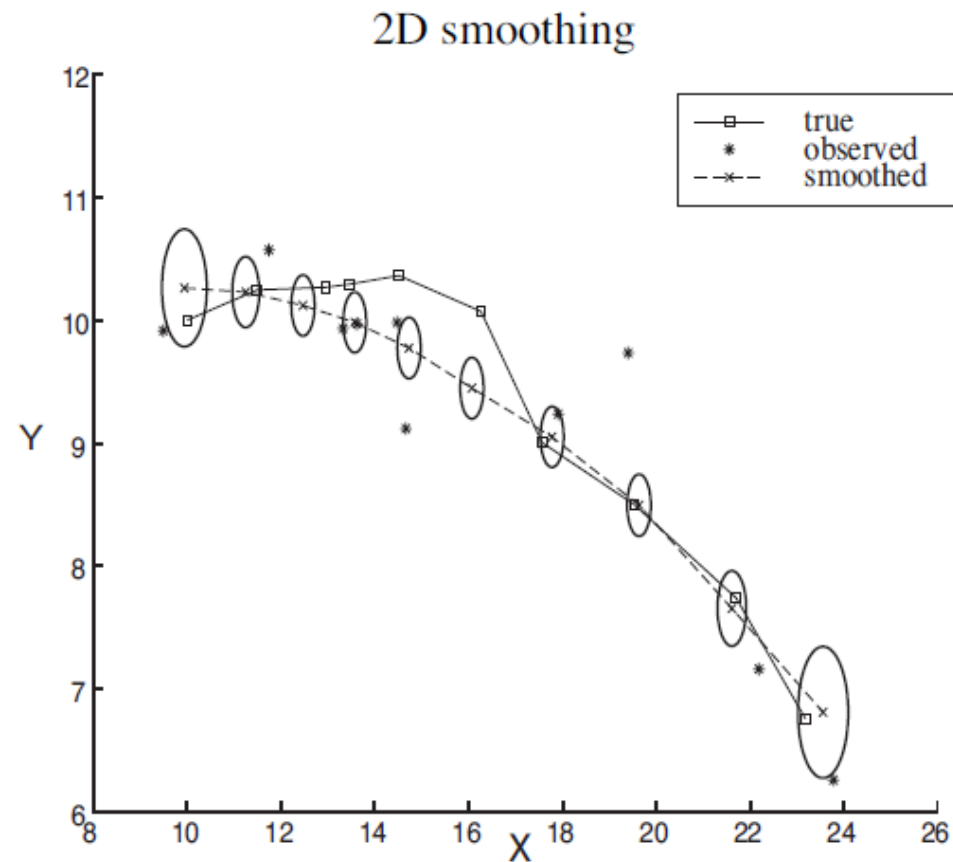
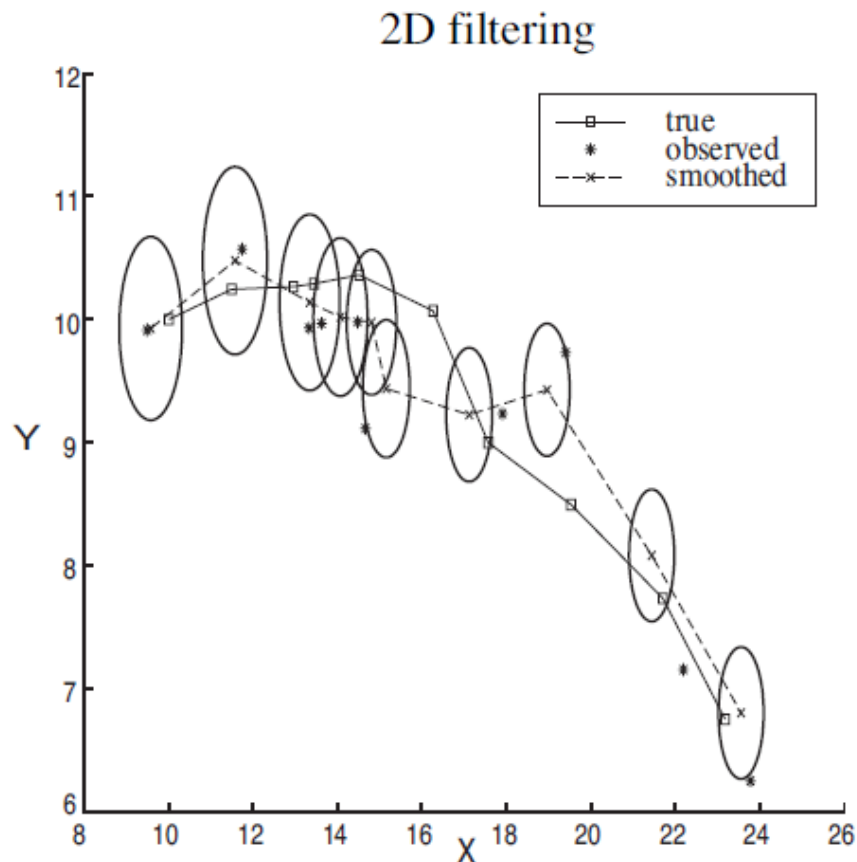


$$\begin{aligned} \mathbf{P}(\mathbf{X}_{t+1} | \mathbf{e}_{1:t+1}) &= \mathbf{P}(\mathbf{X}_{t+1} | \mathbf{e}_{1:t}, \mathbf{e}_{t+1}) \\ &= \alpha \mathbf{P}(\mathbf{e}_{t+1} | \mathbf{X}_{t+1}, \mathbf{e}_{1:t}) \mathbf{P}(\mathbf{X}_{t+1} | \mathbf{e}_{1:t}) \\ &= \alpha \mathbf{P}(\mathbf{e}_{t+1} | \mathbf{X}_{t+1}) \mathbf{P}(\mathbf{X}_{t+1} | \mathbf{e}_{1:t}) \end{aligned}$$

- Start with a Gaussian prior $f_{1:0} = P(X_0) = N(\mu_0, \Sigma_0)$, filtering with a linear Gaussian model produces a Gaussian state distribution for all time.

Kalman Filtering

- Example:
 - A robot moving in a plane with constant but noisy velocity, where only position is observed.



Kalman Filtering


- Linear Gaussian assumption is too strong in real application
- Extensions:
 - Extended Kalman filter
 - Unscented Kalman filter

Inference in DBNs

- Exact inference
 - Variable Elimination
- Approximate inference
 - Particle filter

Particle Filtering

- Problem for filtering:
 - The state distributions grow without bound over time. That is, $|X|$ may be too big to store.
- Solution: **approximate inference**
 - Use a set of samples as the forward message, i.e. use the samples as approximate representation of the current state distribution
 - Focus the set of samples on the high-probability regions of the state space, throw away low-probability samples.



Particle Filtering


- Track samples of X , not all values
- Samples are called **particles**
- $P(x)$ approximated by number of particles with value x
- More particles, more accuracy



Particle Filtering

- **Step 1: Transition (time update)**
 - Each particle is moved by sampling its next position from the transition model

$$x_{t+1} = \text{sample}(P(X_{t+1}|x_t)) \quad \leftarrow \text{Forward sampling}$$




Particle Filtering

- **Step 2: Weighting (Measurement update)**
 - Fix sample observation (evidence)
 - Down weight samples based on the evidence ← Likelihood weighting

$$w(x_{t+1}) = P(e_{t+1}|x_{t+1}))$$


- The probabilities don't sum to one since all have been down weighted (in fact they now sum to (N times) an approximation of $P(e)$)



Particle Filtering

■ Step 3: Resampling

- Rather than tracking weighted samples, we resample from our weighted sample distribution
- This is equivalent to renormalizing the distribution




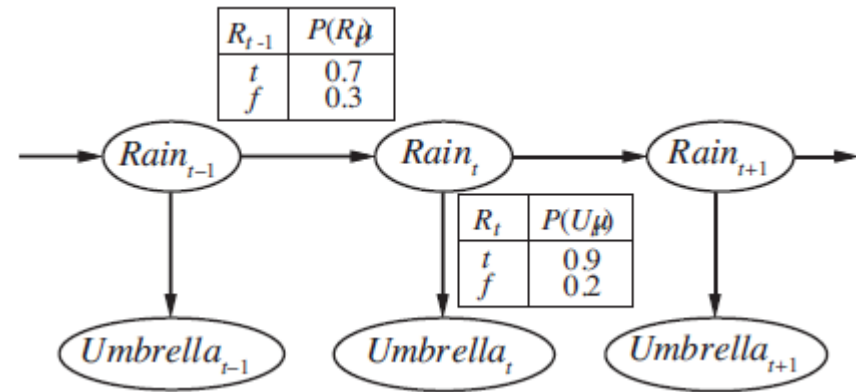
Particle Filtering

```
function PARTICLE-FILTERING(e,  $N$ ,  $dbn$ ) returns a set of samples for the next time step
inputs: e, the new incoming evidence
            $N$ , the number of samples to be maintained
            $dbn$ , a DBN with prior  $\mathbf{P}(\mathbf{X}_0)$ , transition model  $\mathbf{P}(\mathbf{X}_1|\mathbf{X}_0)$ , sensor model  $\mathbf{P}(\mathbf{E}_1|\mathbf{X}_1)$ 
persistent:  $S$ , a vector of samples of size  $N$ , initially generated from  $\mathbf{P}(\mathbf{X}_0)$ 
local variables:  $W$ , a vector of weights of size  $N$ 


for  $i = 1$  to  $N$  do
     $S[i] \leftarrow$  sample from  $\mathbf{P}(\mathbf{X}_1 | \mathbf{X}_0 = S[i])$     /* step 1 */
     $W[i] \leftarrow \mathbf{P}(\mathbf{e} | \mathbf{X}_1 = S[i])$             /* step 2 */
     $S \leftarrow$  WEIGHTED-SAMPLE-WITH-REPLACEMENT( $N, S, W$ )    /* step 3 */
return  $S$ 
```

Particle Filtering

- Example:



(a) Propagate




(b) Weight

(c) Resample

Particle Filtering

- Example: Robot Localization




<http://www.cs.washington.edu/research/rse-lab/projects/mcl>

D. Fox, S. Thrun, W. Burgard, and F. Dellaert, Particle Filters for Mobile Robot Localization

Particle Filtering

- Example: Map-Based Probabilistic Visual Self-Localization

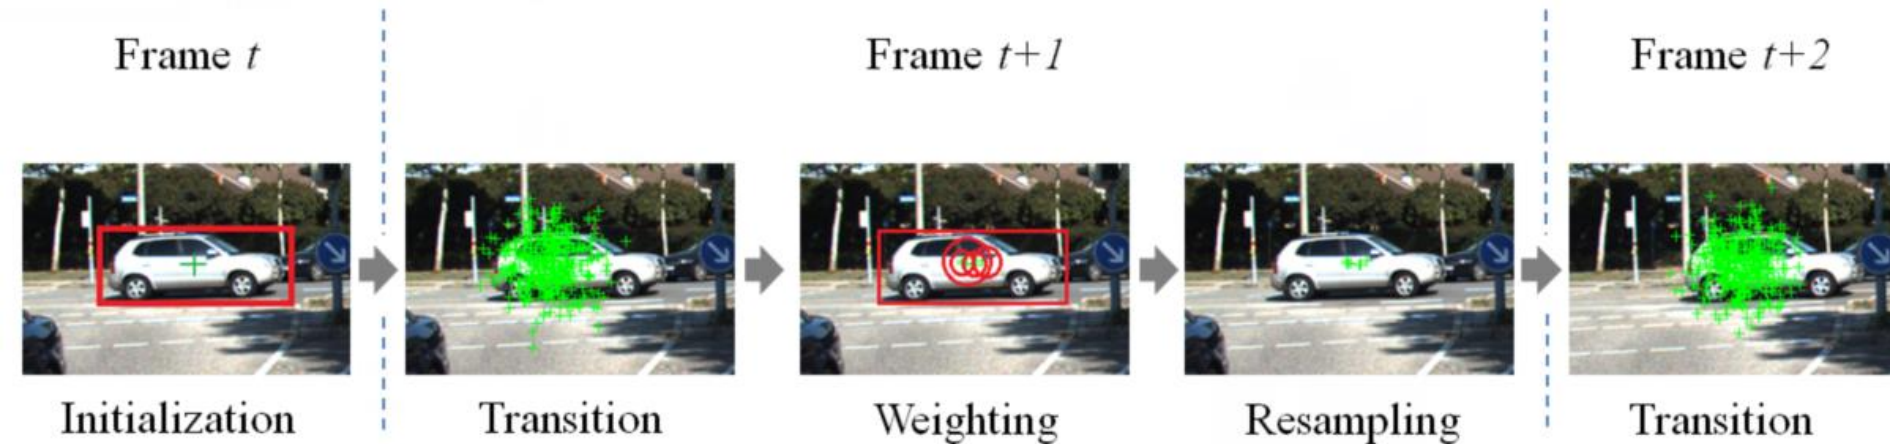


<http://www.cs.toronto.edu/~mbrubake/projects/map/>

M. A. Brubaker, A. Geiger and R. Urtasun, Map-Based Probabilistic Visual Self-Localization, TPAMI 2016

Particle Filtering

- Example: Object Tracking

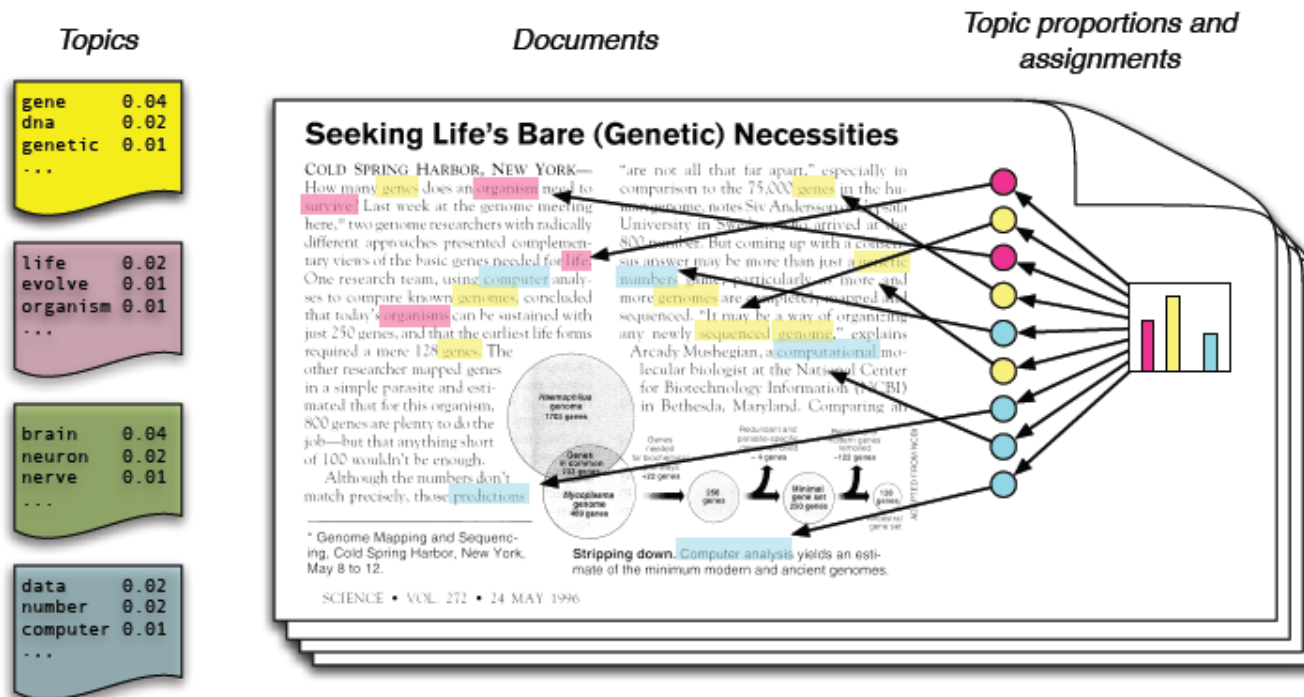


[2] S. Song, Z. Xiang, and J. Liu, Object tracking with 3D LIDAR via multi-task sparse learning, ICMA 2015.

[1] X. Mei and H. Ling, Robust Visual Tracking using l1 Minimization, ICCV 2009


Latent Dirichlet Allocation

- Example: Information retrieval
 - Each topic is a distribution over words
 - Each document is a random mixture of topics
 - Each word is draw from one of those topics



Latent Dirichlet Allocation

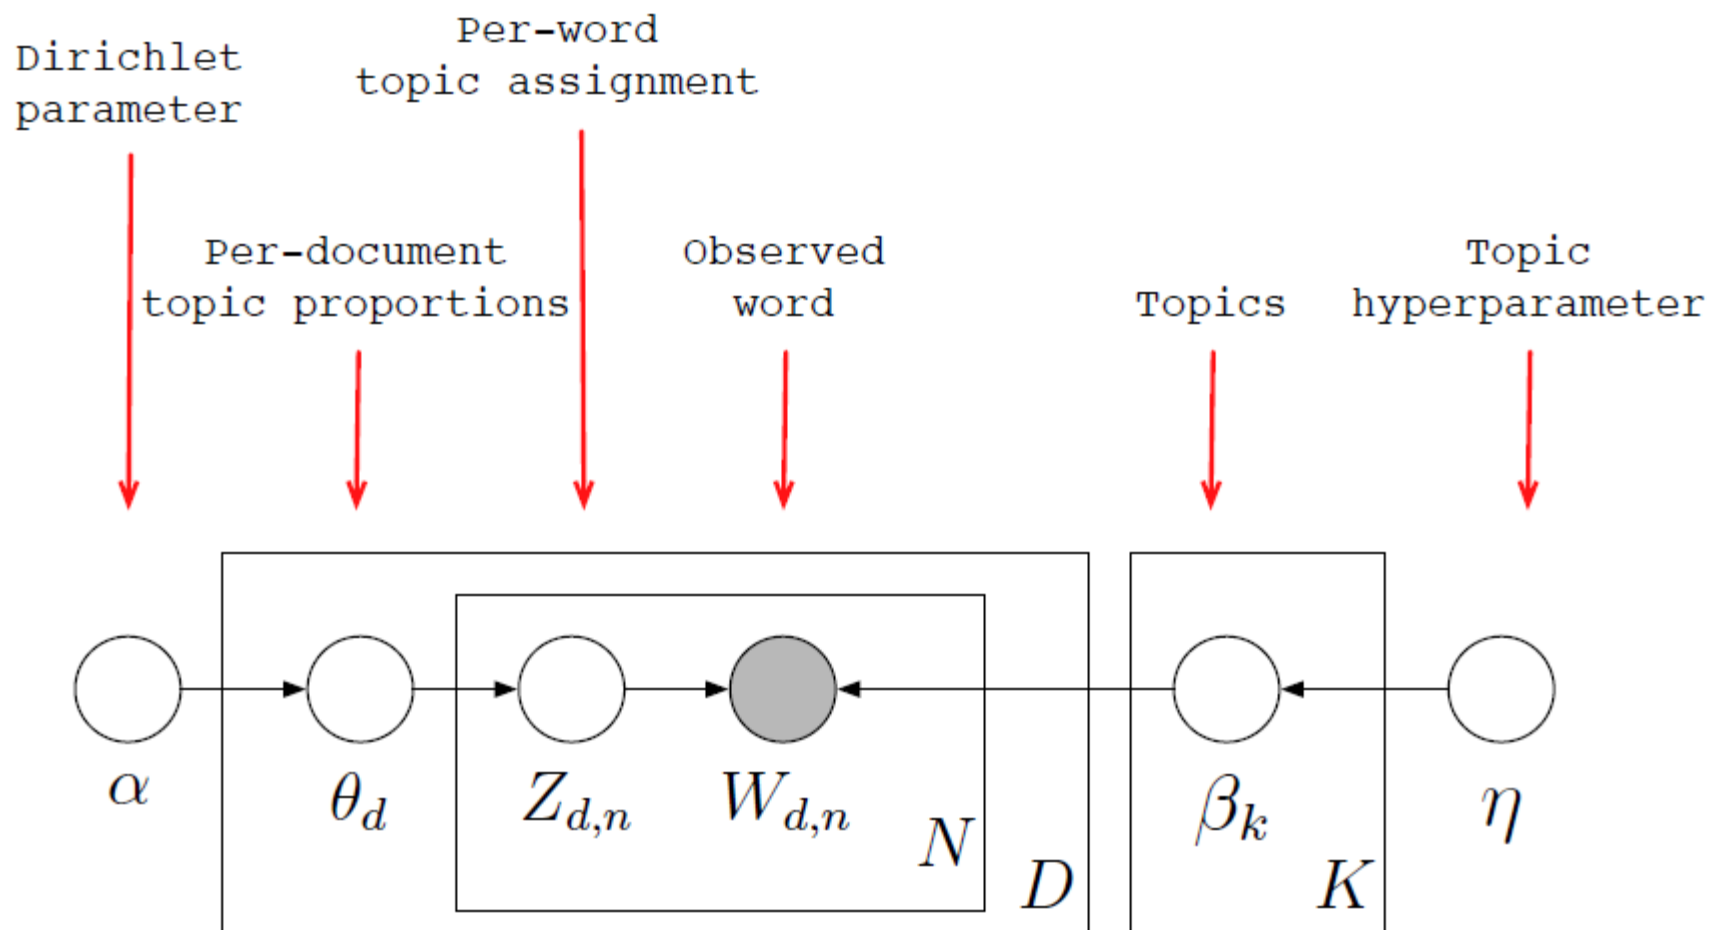
- Example: **Image Parsing**
 - Each topic is a distribution over visual words
 - Each image is a random mixture of topics
 - Each visual word is draw from one of those topics




L. Li, R. Socher, and L. Fei-Fei, Towards Total Scene Understanding: Classification, Annotation and Segmentation in an Automatic Framework, CVPR 2009

Latent Dirichlet Allocation

- LDA Modeling



Latent Dirichlet Allocation




- Joint probability

$$\begin{aligned} &P(Z, W, \theta, \beta, \alpha, \eta) \\ &= P(\alpha)P(\eta) \prod_{d=1}^D P(\theta_d | \alpha) \prod_{n=1}^N P(Z_{d,n} | \theta_d) \prod_{k=1}^K P(\beta_k | \eta) P(W_{d,n} | \beta_k, Z_{d,n}) \end{aligned}$$

Latent Dirichlet Allocation

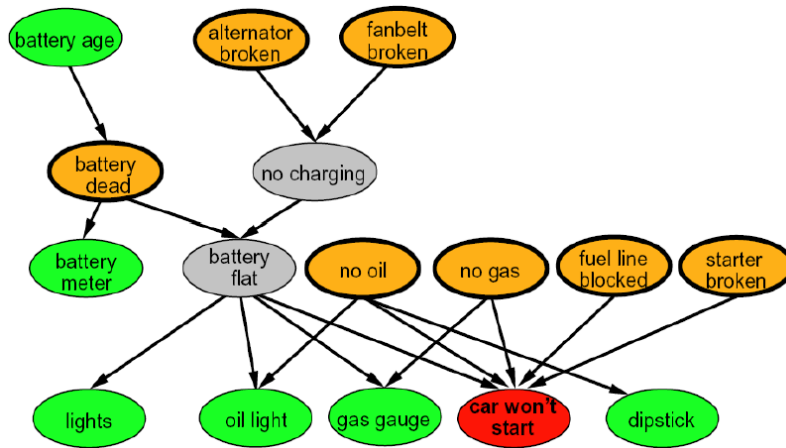
- Inference:




- Maximize the posterior probability $P(Z, \theta, \beta | W, \alpha, \eta)$
- Collapsed Gibbs sampling

Bayesian Network Applications

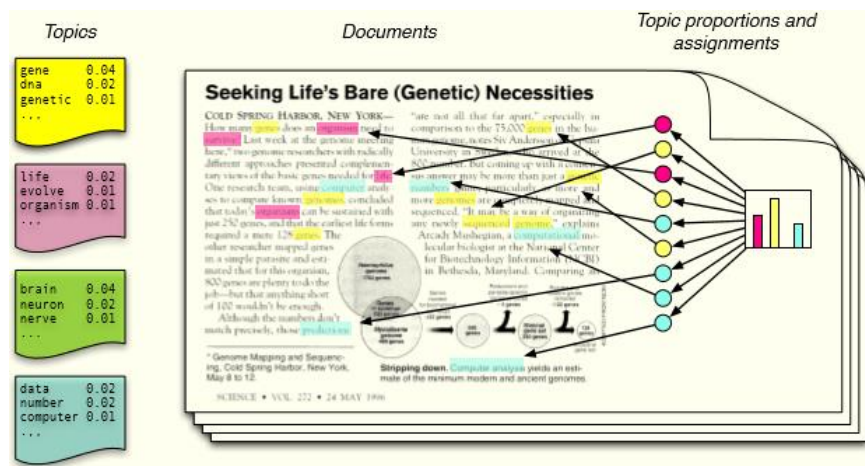
■ Car Diagnosis



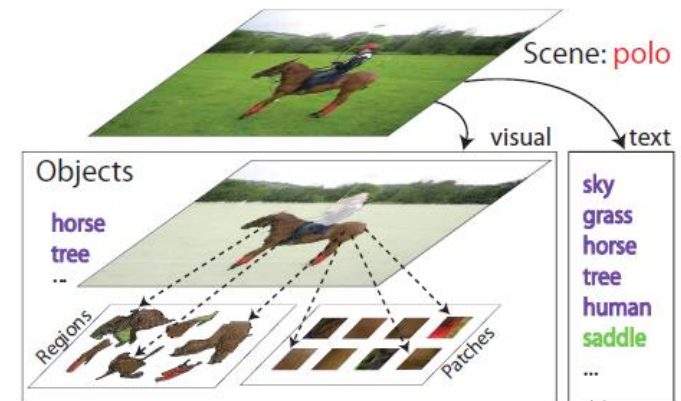
■ Robot Localization



■ Information Retrieval



■ Scene Parsing



Assignments

- Reading assignment:
 - Ch. 15
- Homework 4
- Project 2