

Artificial Intelligence

Lecture 13: Markov Decision Processes

Xiaojin Gong

2021-06-21

Outline

- Markov Decision Processes
 - Problem Formulation
 - Value Iteration
 - Policy Iteration

Review

Problems solved by searching



Probabilistic graphical models

- Problem formulation
- Solving by searching
 - Deterministic environment


- Representation
- Inference
- Learning

Uncertainty

Planning in
deterministic environments



Planning in
stochastic environments




state s , action a



state $\text{Successor}(s, a)$

state s , action a




state s_1

state s_2

state s_3



Applications



- **Cleaning robot**: decide where to move, but actuators may fail, hit obstacles, etc.
- **Network routing**: decide which server to go, but server may crash, etc.
- **Agriculture**: decide what to plant, but don't know weather and thus crop yield
- Elevator scheduling
- Aircraft navigation
- Resource allocation


Example: Robot Navigation

- The problem:
 - The agent lives in a grid
 - Walls block the agent's path
- Noisy movement:
 - 80% of the time, the action North takes the agent North (if there is no wall there)
 - 10% of the time, North takes the agent West; 10% East
 - If there is a wall in the direction the agent would have been taken, the agent stays put
- The agent receives **rewards** each time step
 - Small “living” reward each step (can be negative)
 - Big rewards come at the end (good or bad)
- Goal: maximize the **utility** (sum of rewards)



Markov Decision Processes

- An **Markov Decision Process (MDP)** is defined by:
 - A set of states $s \in S$
 - A set of actions $a \in A$
 - A **transition function** $T(s, a, s')$
 - Probability from s leads to s' , i.e. $P(s'|s, a)$
 - Also called the model or the dynamics
 - A **reward function** $R(s, a, s')$
 - Sometimes just $R(s)$
 - A **start state**
 - Maybe a **terminal state**
- **Solution:**
 - A **policy** π gives an action for each state
- MDPs are **non-deterministic search problems**



Markov Decision Processes

- “Markov” generally means that given the present state, the future and the past are independent
- For Markov decision processes, “Markov” means action outcomes depend only on the current state

$$\begin{aligned} &P(S_{t+1} = s' | S_t = s_t, A_t = a_t, S_{t-1} = s_{t-1}, A_{t-1}, \dots, S_0 = s_0) \\ &= P(S_{t+1} = s' | S_t = s_t, A_t = a_t) \end{aligned}$$



- This is just like search, where the successor function could only depend on the current state (not the history)

Markov Decision Processes

- A sequential decision problem for a fully observable, stochastic environment with a Markovian transition model and additive rewards is called a Markov Decision Process.




Policies

- In deterministic single-agent search problems, we wanted an optimal **plan**, or sequence of actions, from start to a goal
- For MDPs, we want an **optimal policy** $\pi^*: S \rightarrow A$
 - A policy π gives an action for each state
 - An optimal policy is one that **maximizes expected utility** if followed
 - An explicit policy defines a reflex agent




Optimal Policies


- Balancing of risk and reward




$$R(s) < -1.6284$$



$$-0.4278 < R(s) < -0.0850$$




$$-0.0221 < R(s) < 0$$



$$R(s) > 0$$

Utilities of Sequences

- Finite horizon: (similar to depth-limited search)
 - Terminate episodes after a fixed T steps (e.g. life)
 - Gives **nonstationary policies** (π depends on time left)



$R(s) > 0$

- Infinite horizon:
 - No fixed time limit
 - The optimal policy is **stationary**

Utilities of Sequences

- Theorem: if we assume stationary preferences:

$$\begin{array}{c} [s_1, s_2, \dots] \succ [s'_1, s'_2, \dots] \\ \Updownarrow \\ [s_0, s_1, s_2, \dots] \succ [s_0, s'_1, s'_2, \dots] \end{array}$$

- Then: there are only two ways to define utilities

- Additive rewards:

$$U_h([s_0, s_1, s_2, \dots]) = R(s_0) + R(s_1) + R(s_2) + \dots$$

- Discounted rewards:

$$U_h([s_0, s_1, s_2, \dots]) = R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots$$


Utilities of Sequences

- How to discount?
 - Each time we descend a level, we multiply in the discount once
- Why discount?
 - Sooner rewards probably do have higher utility than later rewards
 - Also helps our algorithms converge
- Example: discount of 0.5
 - $U([1,2,3]) = 1*1 + 0.5*2 + 0.25*3$
 - $U([1,2,3]) < U([3,2,1])$
- With discounted rewards, the utility of an infinite sequence is finite.

$$U_h([s_0, s_1, s_2, \dots]) = \sum_{t=0}^{\infty} \gamma^t R(s_t) \leq \sum_{t=0}^{\infty} \gamma^t R_{\max} = R_{\max} / (1 - \gamma)$$


MDP Search Trees

- Each MDP state gives an expectimax-like search tree




Expectimax Search Tree

- Expectimax search: compute the average score under optimal play
 - **Max nodes** as in minimax search
 - **Chance nodes** are like min nodes but the outcome is uncertain
 - Calculate their **expected utilities**, i.e. take weighted average (expectation) of children



Optimal Quantities

- The value (utility) of a state s :
 - $V^*(s)$ = expected utility starting in s and acting optimally
- The value (utility) of a q-state (s,a) :
 - $Q^*(s, a)$ = expected utility starting out having taken action a from state s and (thereafter) acting optimally
- The optimal policy:
 - $\pi^*(s)$ = optimal action from state s
- Optimal values define optimal policies!



Optimal Quantities

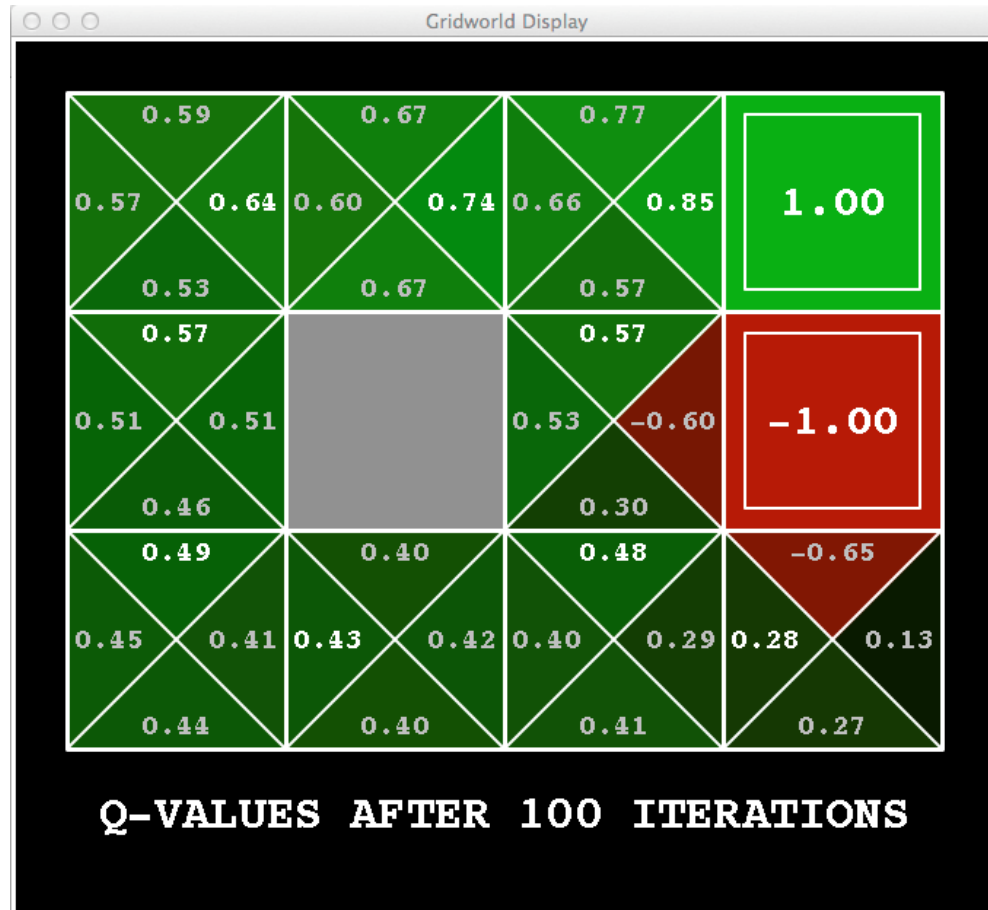
- Example: Gridworld V Values



Noise = 0.2
Discount = 0.9
Living reward = 0

Optimal Quantities


- Example: Gridworld Q Values



Noise = 0.2
Discount = 0.9
Living reward = 0

Optimal Quantities

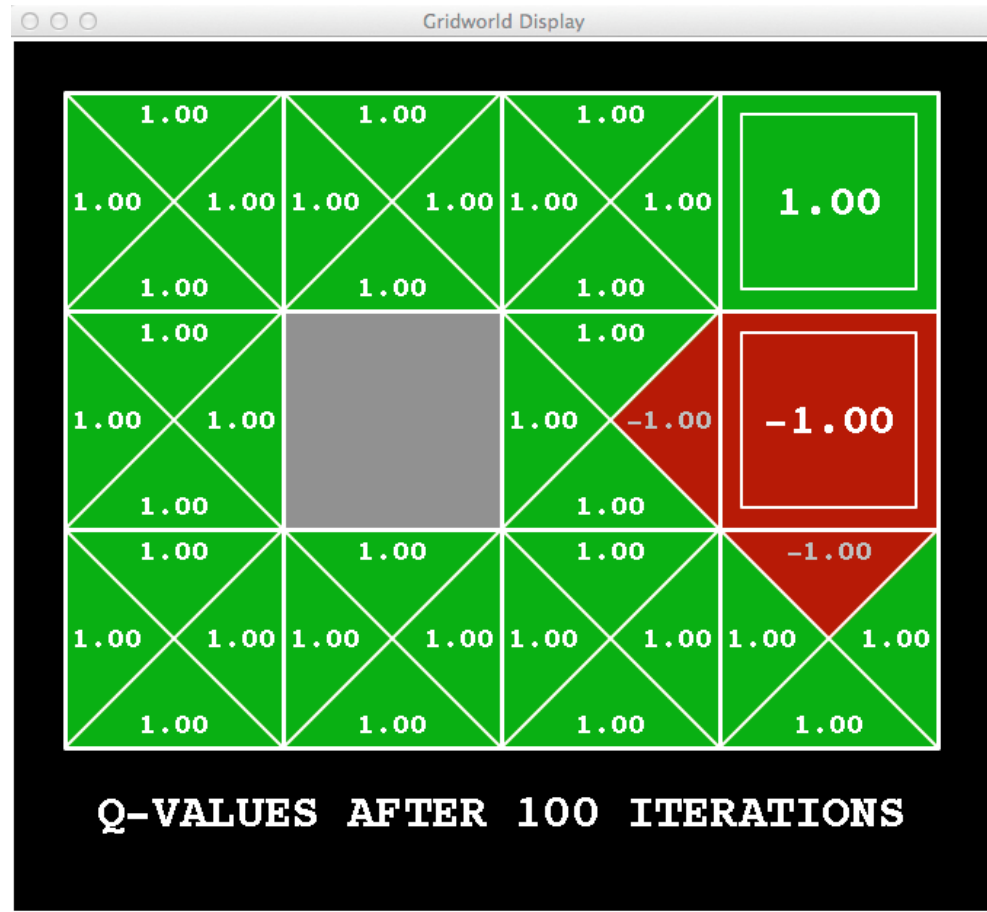
- Example: Gridworld V Values



Noise = 0
Discount = 1
Living reward = 0

Optimal Quantities

- Example: Gridworld Q Values



Noise = 0
Discount = 1
Living reward = 0


Recap: Defining MDPs

- Markov decision processes:

- States S
- Actions A
- Transitions $P(s'/s, a)$ (or $T(s, a, s')$)
- Rewards $R(s, a, s')$ (and discount γ)
- Start state s_0

- Quantities:

- Policy = map of states to actions
- Utility = sum of discounted rewards
- Values = expected future utility from a state (max node)
- Q-Values = expected future utility from a q-state (chance node)




Values of States

- Fundamental operation: compute the (expectimax) value of a state
 - Expected utility under optimal action
 - Average sum of (discounted) rewards
- Recursive definition of value:

$$V^*(s) = \max_a Q^*(s, a)$$

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$



The Bellman Equation


选择期望最大的

$$U(s) = \overset{\text{reward}}{R(s)} + \gamma \max_{a \in A(s)} \sum_{s'} P(s' | s, a) U(s')$$

- Optimal rewards = maximize over first action and then follow optimal policy


- Recursive definition of value:

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$




MDP Search Trees

- Each MDP state gives an expectimax-like search tree



MDP Search Trees

- Problems:
 - This tree is usually infinite
 - Same states appear over and over
 - We would search once per state
- Idea: **Value iteration**
 - Compute optimal values for all states all at once using successive approximations
 - Will be a bottom-up dynamic program similar in cost to memoization
 - Do all planning offline, no replanning needed!




Value Iteration

- Start with $V_0(s) = 0$: no time steps left means an expected reward sum of zero
- Given vector of $V_k(s)$ values, do one ply of expectimax from each state:

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

- Repeat until convergence

- Complexity of each iteration: $O(S^2A)$



The Value Iteration Algorithm

function VALUE-ITERATION(mdp, ϵ) **returns** a utility function

inputs: mdp , an MDP with states S , actions $A(s)$, transition model $P(s' | s, a)$, rewards $R(s)$, discount γ

ϵ , the maximum error allowed in the utility of any state

local variables: U, U' , vectors of utilities for states in S , initially zero

δ , the maximum change in the utility of any state in an iteration

repeat 控制算法收敛的条件

$U \leftarrow U'; \delta \leftarrow 0$

for each state s **in** S **do**

$$U'[s] \leftarrow R(s) + \gamma \max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s']$$

if $|U'[s] - U[s]| > \delta$ **then** $\delta \leftarrow |U'[s] - U[s]|$

until $\delta < \epsilon(1 - \gamma)/\gamma$

return U

- **Theorem:** will converge to unique optimal values
 - Basic idea: approximations get refined towards optimal values
 - Policy may converge long before values do

The Value Iteration Algorithm

$$U(s) = R(s) + \gamma \max_{a \in A(s)} \sum_{s'} P(s' | s, a) U(s')$$

80%成功.

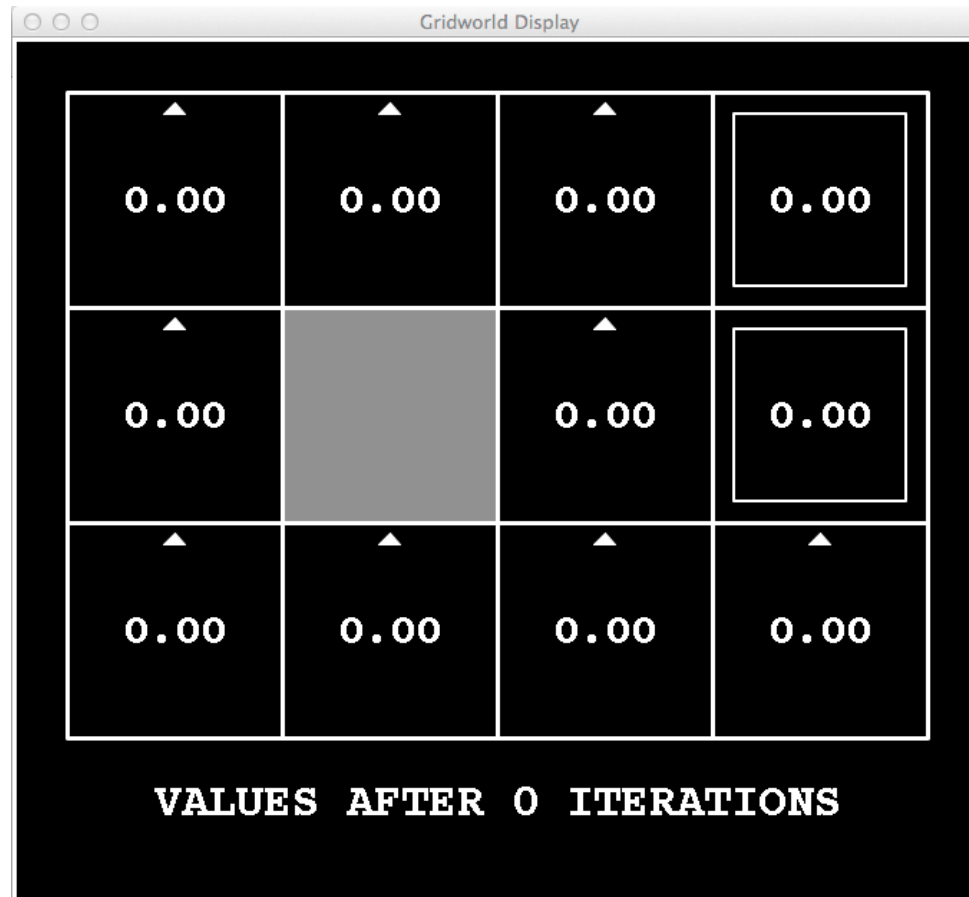
$$U(1, 1) = -0.04 + \gamma \max \left[\begin{array}{l} 0.8U(1, 2) + 0.1U(2, 1) + 0.1U(1, 1), \quad (Up) \\ 0.9U(1, 1) + 0.1U(1, 2), \quad (Left) \\ 0.9U(1, 1) + 0.1U(2, 1), \quad (Down) \\ 0.8U(2, 1) + 0.1U(1, 2) + 0.1U(1, 1) \end{array} \right]. \quad (Right)$$

3	0.812	0.868	0.918	<div style="border: 1px solid black; padding: 2px;">+ 1</div>
2	0.762		0.660	<div style="border: 1px solid black; padding: 2px;">- 1</div>
1	0.705	0.655	0.611	0.388
	1	2	3	4

$\gamma = 1$ and $R(s) = -0.04$ for nonterminal states.

Value Iteration


$k=0$



Noise = 0.2
Discount = 0.9
Living reward = 0

Value Iteration

$k=1$



Noise = 0.2
Discount = 0.9
Living reward = 0

Value Iteration

k=2

最早更新到

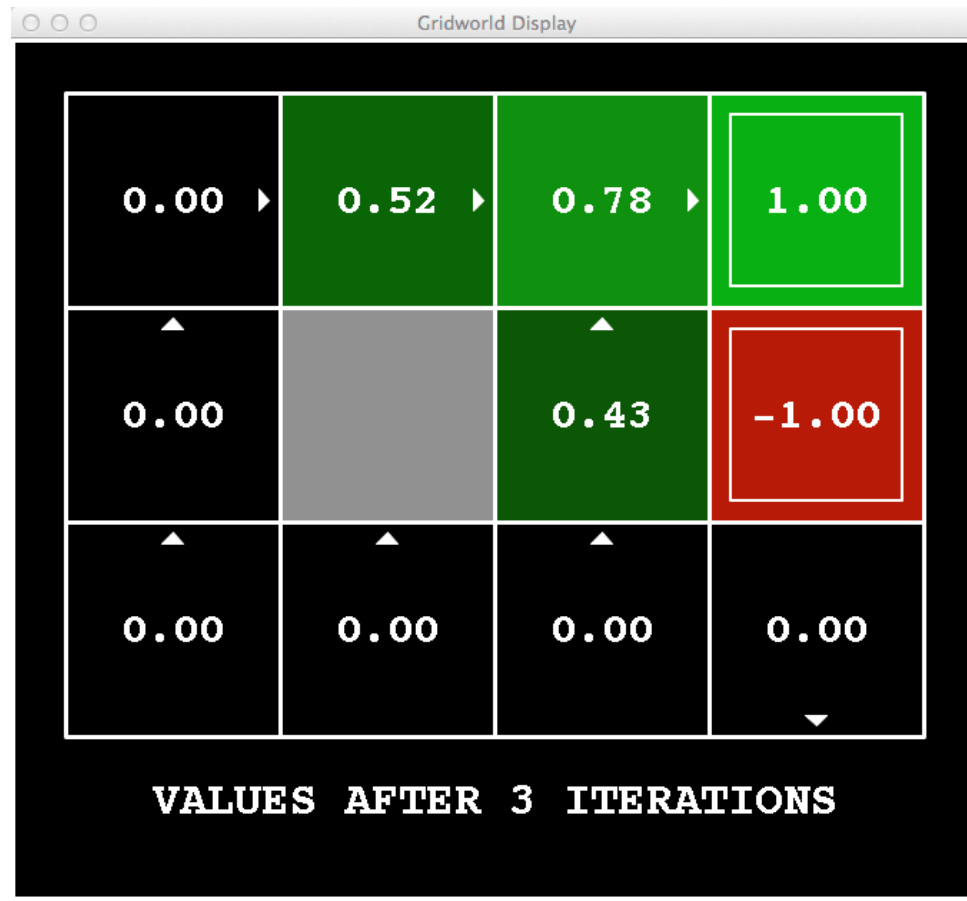
0.9(0.8 , 0.1)



Noise = 0.2
Discount = 0.9
Living reward = 0

Value Iteration

$k=3$



Noise = 0.2
Discount = 0.9
Living reward = 0

Value Iteration

$$k=4$$


Noise = 0.2
Discount = 0.9
Living reward = 0

Value Iteration

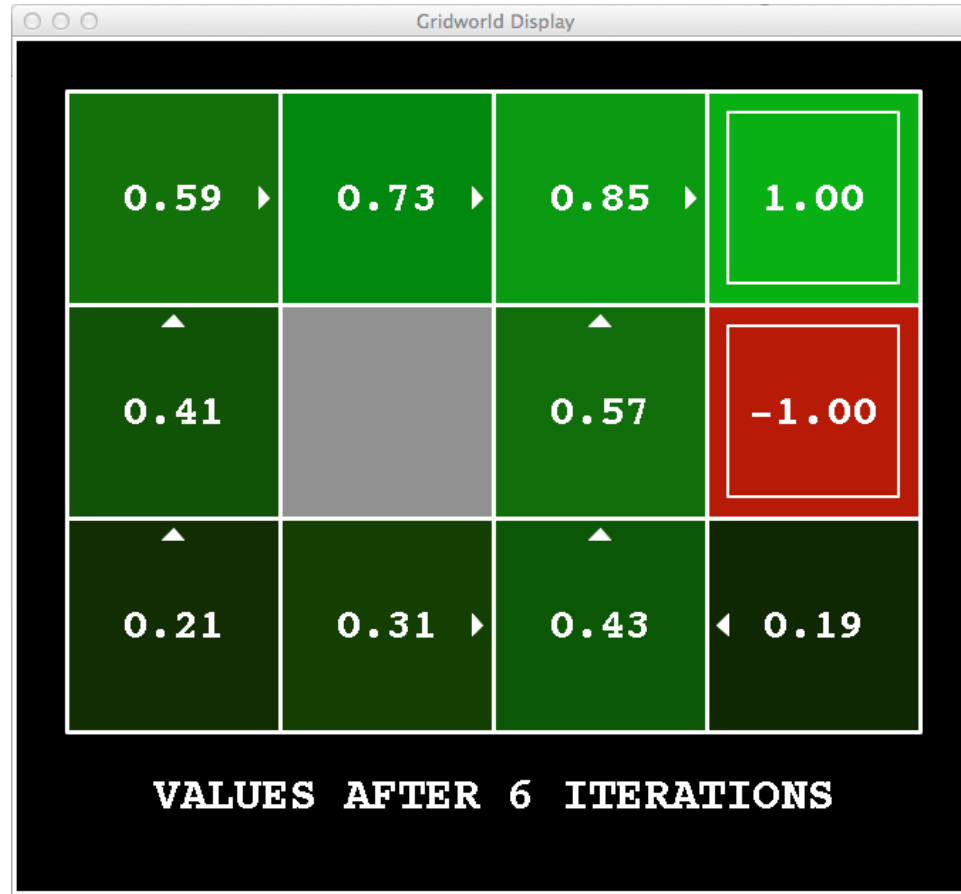
$k=5$



Noise = 0.2
Discount = 0.9
Living reward = 0

Value Iteration

k=6



Noise = 0.2
Discount = 0.9
Living reward = 0

Value Iteration

$k=7$



Noise = 0.2
Discount = 0.9
Living reward = 0

Value Iteration

$k=8$



Noise = 0.2
Discount = 0.9
Living reward = 0

Value Iteration

k=9



Noise = 0.2
Discount = 0.9
Living reward = 0

Value Iteration

k=10



Noise = 0.2
Discount = 0.9
Living reward = 0

Value Iteration

k=11



Noise = 0.2
Discount = 0.9
Living reward = 0

Value Iteration


k=12



Noise = 0.2
Discount = 0.9
Living reward = 0

Value Iteration

k=100



Noise = 0.2
Discount = 0.9
Living reward = 0

Problems with Value Iteration

- Value iteration repeats the Bellman updates:

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

- Problems:
 - It's slow – $O(S^2A)$ per iteration
 - The “max” at each state rarely changes
 - The policy often converges long before the values

Policy Iteration

- The policy iteration algorithm alternates the following two steps, beginning from some initial policy π_0 :
 - Policy evaluation:
 - calculate utilities for some fixed policy (not optimal utilities!) until convergence
 - Policy improvement:
 - update policy using one-step look-ahead with resulting converged (but not optimal!) utilities as future values
- Repeat steps until policy converges

The Policy Iteration Algorithm

function POLICY-ITERATION(*mdp*) **returns** a policy

inputs: *mdp*, an MDP with states S , actions $A(s)$, transition model $P(s' | s, a)$

local variables: U , a vector of utilities for states in S , initially zero

π , a policy vector indexed by state, initially random

policy向量就是每个状态应该采取什么动作

repeat

$U \leftarrow \text{POLICY-EVALUATION}(\pi, U, mdp)$

$unchanged? \leftarrow \text{true}$

for each state s **in** S **do**

if $\max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s'] > \sum_{s'} P(s' | s, \pi[s]) U[s']$ **then do**

$\pi[s] \leftarrow \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s']$

$unchanged? \leftarrow \text{false}$


until $unchanged?$

return π


Policy Evaluation

告诉你action, 看q state

Do the optimal action



Do what π says to do



- Expectimax trees max over all actions to compute the optimal values
- If we fixed some policy $\pi(s)$, then the tree would be simpler – only one action per state
 - The tree's value would depend on which policy we fixed


Policy Evaluation

- Define the utility of a state s , under a fixed policy π :

$V^\pi(s)$ = expected total discounted rewards starting in s and following

- Recursive relation (one-step look-ahead / Bellman equation):

$$V^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^\pi(s')]$$



Policy Extraction

- Computing actions from values

- Let's imagine we have the optimal values $V^*(s)$

- How should we act?
 - It's not obvious!




- We need to do one step

$$\pi^*(s) = \arg \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

Policy Extraction

- Computing actions from values
 - Let's imagine we have the optimal q-values:
- How should we act?
 - Completely trivial to decide!

$$\pi^*(s) = \arg \max_a Q^*(s, a)$$



- Actions are easier to select from q-values than values!

Policy Iteration

- **Evaluation:** For fixed current policy π , find values with policy evaluation

- Iterate until values converge:

$$V_{k+1}^{\pi_i}(s) \leftarrow \sum_{s'} T(s, \pi_i(s), s') [R(s, \pi_i(s), s') + \gamma V_k^{\pi_i}(s')]$$

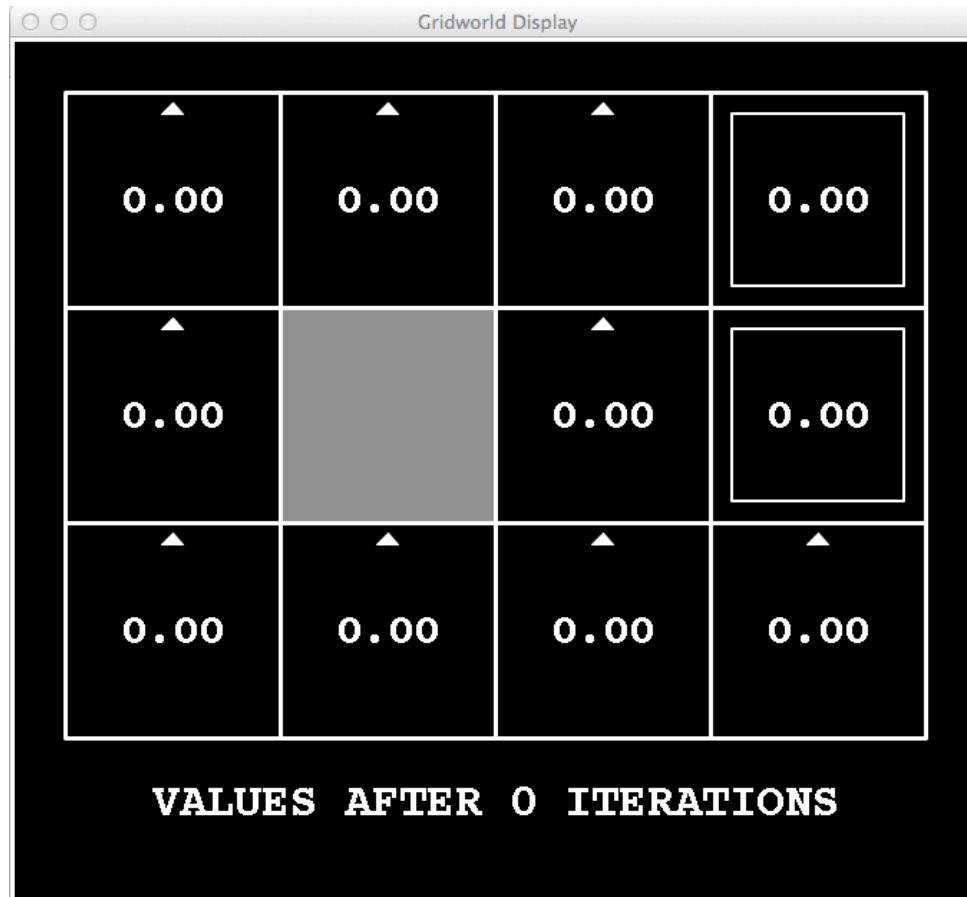
- **Improvement:** For fixed values, get a better policy using policy extraction

- One-step look-ahead

$$\pi_{i+1}(s) = \arg \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^{\pi_i}(s')]$$

Policy Iteration


$k=0$



Noise = 0.2
Discount = 0.9
Living reward = 0

Policy Iteration

$k=1$



Noise = 0.2
Discount = 0.9
Living reward = 0

Policy Iteration

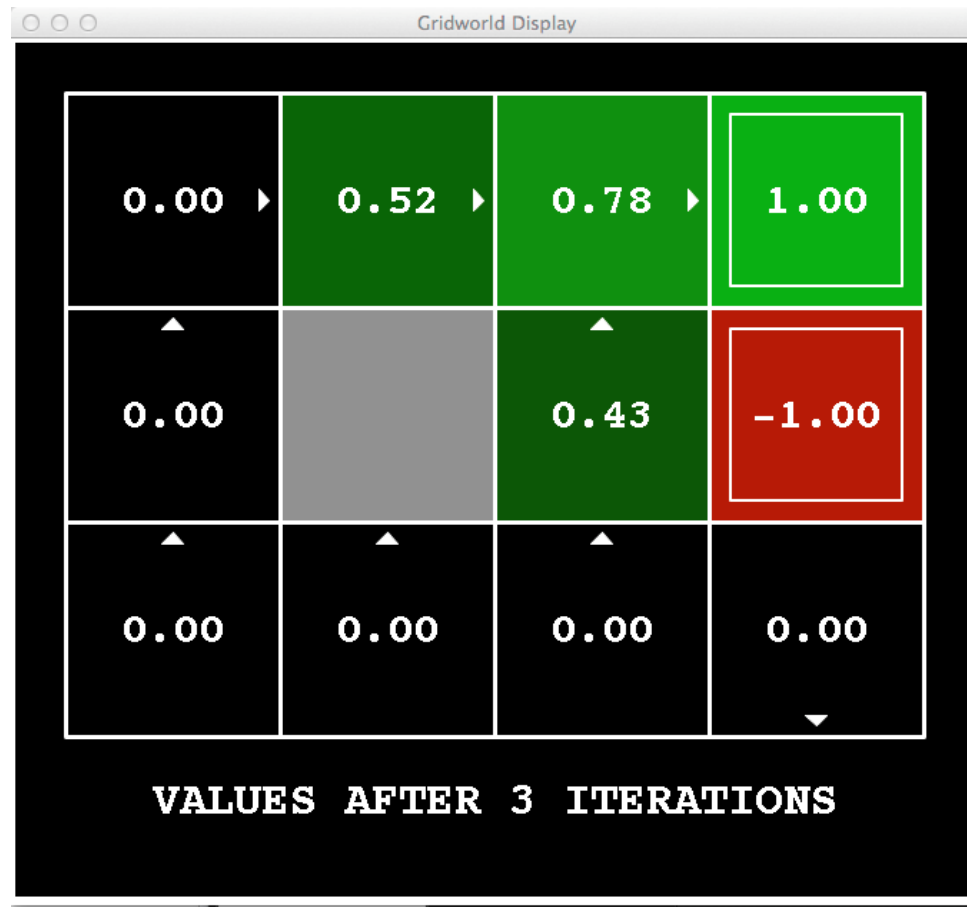
$k=2$



Noise = 0.2
Discount = 0.9
Living reward = 0

Policy Iteration

$k=3$



Noise = 0.2
Discount = 0.9
Living reward = 0

Policy Iteration


k=4



Noise = 0.2
Discount = 0.9
Living reward = 0

Policy Iteration


k=5



Noise = 0.2
Discount = 0.9
Living reward = 0

Policy Iteration

k=6



Noise = 0.2
Discount = 0.9
Living reward = 0

Policy Iteration

$k=7$



Noise = 0.2
Discount = 0.9
Living reward = 0

Policy Iteration

$k=8$



Noise = 0.2
Discount = 0.9
Living reward = 0

Policy Iteration

k=9



Noise = 0.2
Discount = 0.9
Living reward = 0

Policy Iteration


k=10



Noise = 0.2
Discount = 0.9
Living reward = 0

Policy Iteration


k=11



Noise = 0.2
Discount = 0.9
Living reward = 0

Policy Iteration


k=12



Noise = 0.2
Discount = 0.9
Living reward = 0

Policy Iteration

k=100



Noise = 0.2
Discount = 0.9
Living reward = 0

Comparison

- Both value iteration and policy iteration compute the same thing (all optimal values)
- In **value iteration**:
 - Every iteration updates both the values and (implicitly) the policy
 - We don't track the policy, but taking the max over actions implicitly recomputes it
- In **policy iteration**:
 - We do several passes that update utilities with fixed policy (each pass is fast because we consider only one action, not all of them)
 - After the policy is evaluated, a new policy is chosen (slow like a value iteration pass)
 - The new policy will be better (or we're done)

Summary

- So you want to....
 - Compute optimal values: use value iteration or policy iteration
 - Compute values for a particular policy: use policy evaluation
 - Turn your values into a policy: use policy extraction (one-step look ahead)
- These all look the same!
 - They basically are – they are all variations of Bellman updates
 - They all use one-step look ahead expectimax fragments
 - They differ only in whether we plug in a fixed policy or max over actions

Readings

- Artificial Intelligence
 - Chapter 17.1-3
- Final Project
 - Due by July 8, 2021
- Final Exam
 - July 1st, 2021